

# Kernel Threads in xv6

Please **read this entire assignment**, before you start working on the code. This is an especially challenging assignment. You really do not want to wait until the night/day before it is due to start on it. Jeepers, just reading all of ... .. this `__stuff__` is itself enough of a challenge (for a non-engineering student). There may be madness in my method, but there is also method in my madness.

Overall, this assignment is much – much – much more prescriptive<sup>1</sup> than I'd like. I'd prefer it be more descriptive.<sup>2</sup> Following the document does not absolve you of thinking for yourself and filling in places not directly mentioned.

**This lab is May 12<sup>th</sup> by midnight.** Submit a single gzipped `tar` file to **TEACH**. If you don't remember how to create a gzipped `tar` file, you need to learn before you submit this assignment. If your submission is not a gzipped `tar` file, I will not grade your assignment.

There are many parts to this assignment. Many have a lot of steps. Just follow this document like it is a script or recipe and work through all the parts. I recommend you use `#ifdef` sections in your code to make it easier to track where you make changes to the `xv6` source code. I'm sure you've heard that before. I'm sure you have plans to refactor your code after the assignment is due, putting in comments, using mnemonic macros, and using conditional compilation blocks to separate new and old code. Instead, perform that before the due date.

**This assignment is done entirely in the `xv6` environment.**

This programming project is worth **500 points!!!**

## Part 0 – Clone the `xv6-kthreads` directory

I have created a directory from which you can easily begin your coding journey. The directory is called `xv6-kthreads`. It is in the same location where I keep the rest of my `xv6` code. You can clone a copy of that directory using the following command:

```
~chaneyr/Classes/cs444/xv6/xv6-clone.bash -s xv6-kthreads -d xv6-kthreads
```

That will create for you a directory called `xv6-kthreads` which contains all the beginning code for your kernel threads `xv6` adventure. If you prefer to call it `Lab4`, I'm fine with that. Cloning the `xv6-kthread` directory is not a requirement, but it already includes a lot of code to help you get started. It will help you be done sooner.

I have marked many/most of the places in the code where you'll need to make some changes or add code. They are marked with `KTHREADS` in `#ifdef/#endif` blocks. Currently within those blocks are `#error` preprocessor directives. If you enable the `KTHREADS` in the `Makefile`, those will cause the compiler to emit the error message and stop. You'll need to

---

<sup>1</sup> Saying exactly what must happen, especially by giving an instruction or making a rule.

<https://dictionary.cambridge.org/us/dictionary/english/prescriptive>

<sup>2</sup> Presenting observations about the characteristics of someone or something.

<https://www.merriam-webster.com/dictionary/descriptive>

remove the `#error` directives as you develop the code. The beginning code will compile and run. The beginning code includes most of the things we've added for class, such as the `halt` command and the `kdebug` command.

## Part 1 – Add some additional programs – (5 points)

You need to add a couple new user commands/programs into your `xv6-kthreads` directory for some testing. If you are starting with a clone of my `xv6-kthreads` directory, this entire part is already done for you. If not, copy the following programs from my `xv6-kthreads` directory. You can see sample output from the testing programs (`thtst[1-6]`) at end of this document.

<b>memalign.c</b> – this is just a small program to show what steps are taken to assure a block of memory is page aligned.
<b>thtst1.c</b> – the simplest of simple thread tests. It creates a thread, makes sure that memory between the new thread and the main thread is shared, joins with the thread, and exits. This program uses assert statements to validate values in variables for correctness.
<b>thtst2.c</b> – another small simple test of the thread functions, but this one can run with multiple threads, as given from the command line. I've run it with 10 threads. I don't know how much higher it can go under the current limits of <code>xv6</code> . It will simply compute away for a few seconds (usually 30 seconds with 10 threads and 13 seconds with 3 threads, when run in <code>qemu</code> emulating 4 CPUs).
<b>thtst3.c</b> – this is a small version of a multi-threaded matrix multiplication. It generates the same data instead of reading files, it only works on matrices of up to 30x30 (which is the default), it is limited to at most 10 threads, and the output always goes into the same named file ( <code>op.txt</code> ). If you really want to try a larger matrix or more threads, go ahead, but the process memory limit of <code>xv6</code> (as we have it configured) will limit you. Interestingly, most of the run time for this program is spent writing the output file.
<b>thtst4.c</b> – this program tests 2 things. 1) is the parent for a thread correctly established. 2) can a thread join with a thread it did not create. There is nothing complex about this code, it simply creates a few threads, runs for a bit, and joins with them.
<b>thtst5.c</b> – this program tests 1 thing. 1) you should receive a graceful error if you try to join with a non-existent thread.
<b>thtst6.c</b> – this program tests 1 thing. 1) if you pass a non-page aligned pointer to memory to <code>kthread_create()</code> , it should gracefully reject the new thread creation.

Table 1: Some simple testing programs.

Now modify your `Makefile` so that the new programs are compiled and built into the file system when `xv6` starts. As mentioned above, if you cloned my `xv6-kthreads` directory, this is already done. Add the new programs into the `UPROGS` variable in the `Makefile`.

Though these are the current set of test/interesting programs for this project. It is very possible that additional test programs are developed. I will make them available in the `xv6-kthreads` directory. these programs may show “zombie!” when they complete. You can ignore that.

## Part 2 – Add some data members to the `proc` structure – (5 points)

You need to add some data members to the `struct proc` data structure that is in the file `proc.h`. The data members you add are:

<p><b>int oncpu</b> – we will use this to show the CPU on which a thread/process is currently running. Later in the assignment, we will switch from having a single CPU for <code>gemu</code> to having multiple CPUs. The <code>gemu</code> software will allow up to 8 CPUs, but I’ve always just used 4. Switching from a single CPU to multiple CPUs is just a small change to the <code>Makefile</code>. The right default value is for <code>oncpu</code> should -1. When the process/thread is <code>RUNNING</code>, this should be the CPU number, which you will set in <code>scheduler()</code>.</p>
<p><b>ushort is_thread</b> – this is use to indicate that a <i>thingie</i> taking up a slot in the <code>ptable</code> structure is a thread for a process, not the main process itself. Since we want the thread scheduling to be handled by the kernel, we are going to make them entries in the <code>ptable</code>, just as regular processes are. When determining how an entry in the <code>ptable</code> should be handled (by something like the <code>wait()</code> call), this will be very helpful. Only for thingies created through <code>kthread_create()</code> will this be set to <code>TRUE</code>.</p>
<p><b>ushort is_parent</b> – indicates that this process has (or had) threads within it. Like the <code>is_thread</code> data member, this is useful when managing processes. This is <code>TRUE</code> only for the “parent” thread. The parent thread is the initial stream of execution for a multi-threaded process. The parent was created by a call to <code>fork()/exec()</code>. Only see this to <code>TRUE</code>, when a new thread is created by calling <code>kthread_create()</code>.</p>
<p><b>ushort tid</b> – if this <i>thingie</i> is a kernel thread (which is what you are building), this will hold the unique (to the process) identifier for the thread. The <code>tid</code> will be unique for a process, but may not be unique for across all processes (just like <code>PThreads</code>). The default value for non-threads is 0. For a thread created by <code>kthread_create()</code>, it will be the value of the <code>next_tid</code> data member for the parent thread.</p>
<p><b>ushort next_tid</b> – the main thread keeps track of what is the next unique <code>tid</code> to give to a newly created thread. This value should start at 1 and be incremented each time a new thread for that main thread is created. The first thread created, from a call to <code>kthread_create()</code>, for a process will have a <code>tid</code> of 1.</p>
<p><b>int thread_exit_value</b> – if this <i>thingie</i> is a thread, this is will hold the exit value from that thread (set in <code>kthread_exit</code>). This is best set with the actual thread exit value in the call to <code>kthread_exit()</code>. Initialize it to 0 then set it to the actual value in <code>kthread_exit()</code>.</p>
<p><b>ushort thread_count</b> – applies only on the parent thread, it is the count of the number of current threads. For all “child” threads, it should always be 0.</p>

Table 2: New data members for `proc` structure.

Be sure to initialize all these data members in the `allocproc()` function.

### Part 3 – Copy the `benny_thread` code – (5 points)

Copy the `benny_thread.h` and `benny_thread.c` files from my `xv6-kthreads` directory into your development directory. Modify the `Makefile` to build the `benny_thread.o` file from the `benny_thread.c` file. Modify the `Makefile` so that user programs/commands are linked with the `benny_thread.o` object module.



The modification of the `Makefile` to build the `benny_thread.c` module only requires you add `benny_thread.o` into the `ULIB` macro. Doing this will also cause the user programs/commands to link with the `benny_thread.o` file. Not all of them actually need it, but they are fine with it.

Guess what? If you cloned my `xv6-kthreads` directory, this is already done. Otherwise, copy it from my `xv6-kthreads` directory.

**What are the `benny_thread` functions?** An excellent question. The `benny_thread` functions are just a few **user level** functions that make managing the kernel threads a bit easier. The `benny_thread` functions are all user space functions. All, except `benny_thread_tid()`, make kernel space calls to similarly named functions that run in privileged/kernel mode. The `benny_thread` functions are just wrappers that help with the kernel threads; in the same way that `malloc()` is a user level function that makes memory management easier than having to make a bunch of calls to `sbrk()` to handle the heap.

<code>benny_thread_create</code>	<p>Return type: <b>int</b></p> <p>Parameters:</p> <p><b>abt</b> *: the address of a <code>benny_thread_t</code> data type (typedef-ed in <code>benny_thread.h</code>).</p> <p><b>func</b>: A function pointer. The function is a <code>void</code> return and accepts a single parameter, a <code>void *</code>. This should make you think of the <code>start_routine</code> parameter to the function <code>pthread_create</code>.</p> <p><b>arg_ptr</b>: a <code>void *</code> pointer that represents a pointer sized value for the single parameter that is passed to the function <code>func</code> (from above). This should make you think of <code>arg</code> parameter to the function <code>pthread_create</code>.</p> <p>This function is where the memory allocated from the heap that is used as the stack for the thread is performed.</p>
<code>benny_thread_join</code>	<p>Return type: <b>int</b></p> <p>Parameters:</p> <p><b>abt</b>: a <code>benny_thread_t</code> pointer. The <code>tid</code> is passed on to the <code>kthread_exit</code> function.</p>

	<p>This function is where the memory allocated from the heap that was used as the stack for the thread is deallocated.</p> <p>Any thread can join with another thread, except that no thread can join with the main thread (thread 0 for a process). It is important to note that this is a blocking function.</p> <p>This is also where you can pick up the exit value from a thread.</p>
<code>benny_thread_exit</code>	<p>Return type: <b>int</b></p> <p>Parameters:</p> <p><b>exit_value</b>: the exit value for the thread.</p> <p>This is called by a thread when it is complete and ready to terminate. It is, to die for.</p>
<code>benny_thread_bid</code>	<p>Return type: <b>int</b></p> <p>Parameters:</p> <p><b>abt</b>: a <code>benny_thread_t</code> pointer that the <code>benny_thread_tid</code> function will cast back to the <code>benny_thread_s</code> structure and return the <code>tid</code> of the given thread. The <code>benny_thread_t</code> is considered abstract/opaque to functions outside of the <code>benny_thread.c</code> module.</p> <p>This is when a <code>benny_thread</code> wants to know "Who is that?" of another thread.</p>

Table 3: The `benny_thread` functions.

Any code that wants to use the `benny_thread` functions must include the `benny_thread.h` file (as `thtst?.c` test programs do). While it is possible to directly call the `kthread_*` functions from user mode (as the `benny_thread` functions do), it is simpler to use the wrappers. Simplicity is a benny-fit of the functions.

## Part 4 – Stub out the `kthread_*` functions – (5 points)

In the `proc.c` file, stub out the following `kthread_*` functions:

<code>kthread_create</code>	<p>Return type: <b>int</b></p> <p>Parameters:</p> <p><b>func</b>: A function pointer. The function is a <code>void</code> return and accepts a single parameter, a <code>void *</code>. This should make you think of the <code>start_routine</code> parameter to the function <code>pthread_create</code>. It is modeled after that.</p> <p><b>arg_ptr</b>: a <code>void *</code> pointer that represents a pointer sized value for the single parameter that is passed to the function <code>func</code> (from above). This should make you think of <code>arg</code> parameter to the function <code>pthread_create()</code>.</p> <p><b>tstack</b>: a <code>void *</code> pointer to the space that this newly created thread will use as its stack. The <code>tstack</code> pointer must be a page aligned lump</p>
-----------------------------	---


	<p>of memory from user space (NOT kernel space). The <code>tstack</code> must have been allocated before the call to <code>kthread_create()</code> occurred.</p> <p>This is where an actual kernel thread is created. It gets a spot in the <code>ptable</code>, has a kernel stack allocated, has its state set to <code>RUNNABLE</code>, and so much more...</p>
<code>kthread_join</code>	<p>Return type: <b>int</b></p> <p>Parameters:</p> <p><b>tid</b>: an integer that represents the thread identifier (aka <code>tid</code>) for the thread within this process for which the calling thread will join.</p> <p>This is where most of the cleanup for a thread is done. It is important to note that this is a blocking function. If the passed thread is not yet complete (called <code>kthread_exit()</code>), this function will not return until the thread has terminated.</p>
<code>kthread_exit</code>	<p>Return type: <b>void</b></p> <p>Parameters:</p> <p><b>exit_value</b>: an integer that represents the exit status of the thread.</p> <p>A value of 0 generally means a successful termination of the thread. Any value other than zero generally indicates a non-successful termination of the thread. When testing, we can use non-zero values.</p> <p>This is where a thread declares it is done and terminates. However, even though it is done, it must remain in the <code>ptable</code> (as a <code>ZOMBIE</code>) until another thread joins with it. The brains of the thread are removed and it turns in a zombie. Only when another thread joins with it is it removed from the <code>ptable</code> (just as happens with a process).</p> 

Table 4: The `kthread` functions

If you are starting with a clone of my `xv6-kthreads` directory, these functions are already stubbed out in `proc.c`.

You have soooo much already done and it has been so easy. Sorry partner, but the easy part is about to change. It is time to release your innermost wild kernel hacker.

## Part 5 – Implement `kthread_*` Functions – (100 points)

This is where it starts to be challenging and just downright fun. While the following instructions are extensive, **they are not intended to represent everything necessary in the `kthread_*` functions.**

In addition to the `kthread_*` functions, we will modify a couple other functions in this section. The functions will be validated with the test programs.

When writing these functions, I put several blocks of code that used the `proc_kdebug_level` from the `kdebug` function. It makes it easy to enable and disable diagnostics.

### **`kthread_create()`**

This is going to be a lot like the `fork()` function. In fact, starting with a copy of `fork()` is not a bad idea at all. I'm going to assume you did this and use or rename variables from the `fork()` function below.

The first thing the `kthread_create` function must do is to check to make sure the `tstack` pointer is page aligned. Remember the `memalign` program (mentioned in Table 1 on page 2), look at the source code for it. If it is not page aligned, return -1.

Next, call `allocproc()` and assign the value to the `np` variable (which I renamed `newthread`). I also renamed the `pid` variable as `tid`.

The next thing `fork()` does is make a copy of the page table (the call to `copyuvm`, for copy user virtual memory). But, that is for a process (where each process has its own page table). This is a thread, so all you need to do is have the `pgdir` member of `newthread` point to the `pgdir` of `curproc`. This is one of the most important things about a thread, all threads in a process share a single page table.

The size of the thread (the `sz` member) is the same size as `curproc` (since they are the same process). This is actually an issue, but we address how we display the size later in this lab.

The `tf` data member for a `proc` stands for trap frame (it was not my idea to call it that. I would have used a longer data member name, like `ttf`, for "the trap frame"). The `newthread` has a copy of the `curproc` trap frame. Do this (yes the `*` characters are required, that's how it is copied):

```
*newthread ->tf = *curproc->tf;
```

Leave the line that clears the `eax` data member to zero. It seems to have no effect for the threads.

The `eip` register (in the `tf` structure) represents the instruction pointer for the new thread. You should assign `func` to it. This is not touched in the `fork()` function, ~~but it is initialized in the `allocproc()` function. Use the same syntax as is used in `allocproc()`, but use `func` instead of `forkret`.~~

Make sure you set the `is_thread` member for `newthread` to `TRUE`.

There is not a true parent-child relationship between threads, but we need to keep track of which threads are related in a process. So, set the parent of the new thread `newthread` to `curproc`, **UNLESS** `curproc` is itself a thread. If `curproc` is a thread, then the parent of `newthread` is the parent of `curproc`. Use this opportunity to also



set the `is_parent` member in the process (main thread). After assigning the current value of `next_tid` to the `tid` data member of `newthread`, increment the `thread_count` of the parent process (the one running `main()`). When inspecting the `parent` data member in the `proc` structure, all threads should point back to the main thread (as shown in Figure 1). You do not want to have a multi-level hierarchy of relationships. Part of this can be accomplished like this:

```
tid = newthread->tid = newthread->parent->next_tid++;
newthread->parent->thread_count++;
```




Figure 1: Relationship between newly created threads and the main thread (thread 0). All threads created with `kthread_create()`, for a given process, have the same parent, the main thread.

We need to assign the `esp` (extended stack pointer) data member (from the `tf` structure in `newthread`) to the `tstack` that was passed as a parameter. However, the `tstack` variable is the opposite end of the stack (as the stack grows from low address to high address). So, try something like this:

```
newthread->tf->esp = ((int) tstack) + PGSIZE;
```

We are going with the assumption that the size of the stack for each thread is a single page (`PGSIZE`). It would be nice to be able to create threads with different stack sizes, but that is beyond this assignment.

We need to push a value on the stack. Specifically, we need to push the `arg_ptr` variable value onto the stack for the thread function to pick up. This will take 2 statements. See if these make sense (do them whether they make sense or not).

```
newthread->tf->esp -= sizeof(void *);
*((int *) (newthread->tf->esp)) = (int) arg_ptr;
```

The first line decrements the value for the beginning of the stack pointer. The thread function will pull the value from the `esp` register beginning at this location. The second line copies the value from the `arg_ptr` variable on the stack. We have to do some magic C casts to make sure everything is happy.

We have 2 more manipulations of the stack pointer. We need to push the value of the new thread's `tid` onto the thread stack. This is the return value from `kthread_create` that the calling function can pick up. Consider this

```
newthread->tf->esp -= sizeof(tid);
*((int *) (newthread->tf->esp)) = tid;
```

You are getting to be a real pro at this. We are almost there for `kthread_create`.

In `fork()` there is a little loop where the file descriptors from `curproc` are duplicated (using `filedup()`) into `np`. You need to do this for the thread. It is sad to



do this, I'd rather just set a pointer to the same file descriptors. But, the file descriptors are a hard-coded array.

Do the same assignment for `cwd` as is done in `fork()`. As `fork()` does, do the `safestrcpy()`. Acquire the `ptable` lock, set the state of the thread to `RUNNABLE`, and release the lock. Return the `tid`.

Whew... done with `kthread_create()`. As easy as 1, 2, 3, ...  $\infty$ . Gotta watch that last step. It's a doozy.

### **`kthread_join()`**

The `kthread_join()` function is a bit easier than `kthread_create()`, but it has a few subtle requirements. The `kthread_join()` has a lot in common with the `wait()` function. In fact, you will need to make a change to the `wait()` function later. Anything you do in the `kthread_join()` function that makes a change to the `ptable` process array will require a lock on that structure. Make sure that you also unlock it before returning.

Use a variable with a wild name like `return_value` (initialized to `-1`) to hold the return value for the function.

There is a lot more room for creativity in development of this function. I'm going to walk through how I did it, but there are a lot of opportunities for variation.

I used a variable called `thread0` as the return from the call to `myproc()`. Finding the right `thread0` is very important, but let's get ahead of ourselves.

If `tid == 0`, bail and return a `-1`. We cannot join with `tid 0` (i.e. the `thread0` process).

If `thread0's is_thread` data member is `TRUE`, make `thread0 = thread0->parent`. Remember how we created a single level of threads linked to the parent thread? We are using that here.

Now, acquire the lock on the `ptable` and begin to loop through it, using the `p` variable. We are looking for the thread whose `parent` is `thread0` and whose `tid` data member is equal to the `tid` passed into the function. All other thingies, we ignore.

Now that we've found the `tid` that matches the passes `tid` and it's parent is `thread0`, we need to make sure that thread is complete. For that, we enter a tight `while` loop checking to see when its status becomes `ZOMBIE`. Its status will be set to `ZOMBIE` when it calls `kthread_exit()` on itself. During that `while` loop, we need to make sure we allow other process/threads to run and don't hold onto the `ptable` lock. I go into the following loop:

```
while (p->state != ZOMBIE) {
    release(&ptable.lock);
    yield();
    acquire(&ptable.lock)
}
```

Once we've completed that loop, we've found the right thread and it has completed.

We must decrement the thread count for `thread0`.

Free the `kstack` for the thread. This was allocated through a call to `allowproc()` while in `kthread_create()`.

After calling `kfree()`, we just need to clear out the data members from the `ptable` entry pointed to by `p`. Set `kstack` to `NULL`, `pid` to `-1`, `parent` to `NULL`, `name[0]` to `NULL`, `state` to `UNUSED`, `sz` to `0`, and `killed` to `FALSE`. Set the `return_value` local variable to the data member `thread_exit_value`.

**DO NOT CALL `freevm()`.** That would free the virtual memory for the entire process. We do not want to do that here.

Break out of the `for(;;)` loop, release the `ptable` lock and return the `return_value`.

**Okay, let's go make a little change to the `wait()` function.** If the calling thingie is a thread, panic. If the calling thingie has a `thread_count > 0`, panic. It looks like this:

```
if (curproc->is_thread == TRUE) {
    panic("called wait on thread");
}
if (curproc->thread_count > 0) {
    panic("called wait on thread0 with children");
}
```

Luckily, none of the test programs do these nasty things.

Two big functions down (`kthread_create()` and `kthread_join()`); 1 more big function to go.

### **`kthread_exit()`**

The code for `kthread_exit()` is pretty straight forward.

Get a variable called `curproc` (as `exit()` does). If `curproc` data member `is_thread` is `TRUE`, then:

Close all the open files (see `exit()`).

Cleanup the `cwd` (exactly as `exit()` does it with `begin_op` and `end_op`).

Set `killed` to `FALSE`, the thread data member `thread_exit_value` to the passed `exitValue`, `oncpu` to `-1`, and `state` to `ZOMBIE`.

Now, acquire the `ptable` lock and call `sched()` (**not `scheduler()`**). Follow the call to `sched()` with a `panic("kthread_exit")` call. Obviously, it should never get to the call to `panic`. This function does not return, just like you'd not expect a call to `exit()` to not return.



Okay, let's go make a little change to the `exit()` function. This is exactly like the code we added into the `wait()` function. If the calling thingie is a thread, panic. If the calling thingie has a `thread_count > 0`, panic. It looks like this:

```
if (curproc->is_thread == TRUE) {
    panic("called exit on thread");
}
if (curproc->thread_count > 0) {
    panic("called exit on thread0 with children");
}
```

Luckily, none of the test programs do these nasty things.

That's it for `kthread_exit()`. Time to give yourself a big woot woot!

### `kill()`

There is a small change we need to make to the `kill()` function. We need to add a small block of code that makes the exact same check for a thread being killed or if a parent thread with live threads calls `kill()`.

### `scheduler()`

Since we want to know on which CPU a process or thread is running, we need to update that in the `scheduler()` function.

At the top of the function, create a new variable called something like `current_cpu` and assign it the value from `cpuid()`.

When a process/thread is chosen to be scheduled, assign its `oncpu` data member the value of `current_cpu`.

When the scheduled process/thread completes (following the `switchkvm()`), reset the `oncpu` data member back to -1.

### Sundry files and functions

You are adding 3 new system functions. That means you need to go through the rest of the process of modifying the other files to create the system functions. You should be getting used to this by now. Specifically, you'll need to touch the following files (just as you have for other new system functions):

- `defs.h`
- `syscall.h`
- `syscall.c`
- `sysproc.c`
- `usys.S`

## Part 6 – Refresh the `ps/cps()` code – (20 points)

We've added a couple data members to the `proc` structure (Remember Part 2? Seems like ages ago?). We want those to show up when we run the `ps` command.

You need to add the following to the output from the `cps()` function: `oncpu`, `thread_count`, `is_thread`, and `tid`. The header information should be: `"cpu"`, `"thd cnt"`, `"is thrd"`, and `"thrd #"`.

You only show the `oncpu` value when it is  $\geq 0$ . If you've followed the instructions from above, this should be easy. Only a `RUNNING` process/thread should have a value of `oncpu` that is greater than or equal to 0. Do not show the negative values for `oncpu`. A CPU number should not show up more than once in the `ps` listing.

The `"thd cnt"` column should show the count of active threads only then the process is a parent and has still running threads. Otherwise, show a blank.

The `"is thrd"` column shows a 'y' for each kernel thread. Only the active things that were created from `kthread_create()` show up with a 'y'. All others are an empty space.

Finally, the `"thd #"` should only show the `tid` data member value for those things that were created from `kthread_create()`. Anything not created that way are empty.

The easy way to see this data is to run `"tbtst2 6"` and then run `ps` several times to see the output. If you see a "zombie!" after doing this, don't worry, it is only a flaw in their shell. You can see an example of running `ps` with 4 CPUs and the `tbtst2 6` at the end of the document.

## Part 7 – Update and Validate on 4 CPUs – (60 points)

Change the `CPUS` macro in the `Makefile` to 4. It is fine for you to do your development with a single CPU in `qemu`. However, your code will be tested with the value of the `CPUS` macro in the `Makefile` set to 4. I highly recommend you test your code this way. In the `Makefile`, look around line 256. The points for this part are not awarded to simply changing the `Makefile`, but for all the test programs working correctly with 4 CPUs.

## How It will be Graded

When we grade, we will first run the 6 test commands with a single CPU. We will run 1 test program, then exit `qemu` before running the next test program. I know this stuff is hard and we really don't have the 6 months to develop a full test suite, that's why we will exit `qemu` between tests.

We can run `qemu` using a single CPU with the following command:

```
make nox CPUS=1
```

We can run `qemu` using 4 CPUs with the following command:

```
make nox CPUS=4
```

The macro on the command line will override the setting within the `Makefile`.

## Other Tips – Make Sure You Read This

I have to be honest, there are a couple places where I've struggled when writing this code. One of the biggest is where I called `kfree()` in the `kthread_join()`. If you look in the code for `kfree()`, you'll see that it does a `memset()` on the page of kernel memory to be freed.

Doing the `memset()` is an excellent idea, for the reason mentioned in the comment. However, somewhere I must have some boundary conditions messed up. When I run `thtst2` and do the `memset()`, I will usually get “unexpected trap 14 from ...”. This is especially true when I run with more than 1 CPU. If `#ifdef` out (or comment out) the call to `memset()` in `kfree()`, all is fine. Many web searches later, it would seem that I have overrun a buffer and stomped all over an instruction pointer. But, I cannot see where I’ve gone astray. I tracked address back through `kalloc()` to `kfree()` and found nothing. I would really like to know what I am doing wrong. ☹ If you find out what my error is, have pity on me and let me know.

Another place I struggled was the `suptime()` function we wrote as part of the lottery scheduler. The `suptime()` function makes a lock on the `tickslock`, copies the value of `ticks` (a kernel wide global variable), and releases the lock. No muss – no fuss. Yet, when run with more than 1 CPU, the kernel will hang when trying to acquire that lock. I’ve stepped through this code for hours trying to find what I can do to prevent it. I’ve found nothing. NOTHING!!! You can take the path of least resistance and just enable the `DTICKS` macro in the `Makefile` to get around the issue. If you read this only after spending a few hours debugging, I’m sure you learned something. If you did find how to truly fix this, please let me know.



## Submit to TEACH

When you are done with the Lab4, submit your code to TEACH. Remember how we used the command “`make teach`” to produce a tar and gzipped file that you can submit into TEACH? Do that and be done.

## Final note

The labs in this course are intended to give you basic skills. In later labs, we will *assume* that you have mastered the skills introduced in earlier labs. **If you don’t understand, ask questions.**

## Example Output from Test Programs and ps

```
$ thtst1
global before: 10
i before      : 0xF0F0F0F
rez           : 0x0
global after  : 100
i after       : 0xAEAEAEAE
rez           : 3
$ █
```

```
$ thtst2
Starting 4 threads
thtst2.c 64: started thread 1
thtst2.c 64: started$ thread 2
thtst2.c 64: started thread 3
thtst2.c 64: started thread 4
thtst2.c 69: joining with 1
thtst2.c 69: joining with 2
thtst2.c 69: joining with 3
thtst2.c 69: joining with 4
All threads joined
zombie!
$ █
```

```
$ thtst3
num threads 4
thtst3.c 103$ : 4
thtst3.c 110: 4
  created thread 1 0
  created thread 2 1
  created thread 3 2
  created thread 4 3
  join thread 1 0
  join thread 2 1
  join thread 3 2
  join thread 4 3
zombie!
```

```
$ █
```

```
$ thtst4
$
$
$ All threads joined
zombie!
```

```
$ █
```

```
$ thtst5
Starting 4 threads
thtst5.c 65: 1
thtst5.c 65: 2
thtst5.c 65: 3
thtst5.c 65: 4
thtst5.c 73: -1
thtst5.c 77: joining with 1
thtst5.c 77: joining with 2
thtst5.c 77: joining with 3
thtst5.c 77: joining with 4
All threads joined
zombie!
```

```
$ █
```

```
$ thtst6
*** thread stack not page aligned ***
thtst6.c 34: -1
```

```
$ █
```

```
$ ps
pid      ppid    name      state    size      start time      ticks    sched    cpu    thd cnt is thrd thrd #
1         1      init      sleep    12288     2020-04-30 10:43:22    9        43
2         1       sh        sleep    16384     2020-04-30 10:43:22    5        31
6         5      thtst2    runble   77824     2020-04-30 10:43:33   186      188
5         1      thtst2    runble   77824     2020-04-30 10:43:33    6        324      6      y      1
7         5      thtst2    run      77824     2020-04-30 10:43:33   184      186      2      y      2
8         5      thtst2    runble   77824     2020-04-30 10:43:33   184      184      y      3
9         5      thtst2    run      77824     2020-04-30 10:43:33   183      184      1      y      4
10        5      thtst2    run      77824     2020-04-30 10:43:33   182      183      0      y      5
11        5      thtst2    runble   77824     2020-04-30 10:43:33   182      182      y      6
13        2       ps        run      12288     2020-04-30 10:43:35    0         8      3

$ ps
pid      ppid    name      state    size      start time      ticks    sched    cpu    thd cnt is thrd thrd #
1         1      init      sleep    12288     2020-04-30 10:43:22    9        43
2         1       sh        sleep    16384     2020-04-30 10:43:22    6        34
6         5      thtst2    runble   77824     2020-04-30 10:43:33   311      313
5         1      thtst2    runble   77824     2020-04-30 10:43:33    6       680      6      y      1
7         5      thtst2    run      77824     2020-04-30 10:43:33   309      311      3      y      2
8         5      thtst2    runble   77824     2020-04-30 10:43:33   309      309      y      3
9         5      thtst2    runble   77824     2020-04-30 10:43:33   309      309      y      4
10        5      thtst2    run      77824     2020-04-30 10:43:33   307      308      2      y      5
11        5      thtst2    run      77824     2020-04-30 10:43:33   307      308      0      y      6
14        2       ps        run      12288     2020-04-30 10:43:37    0         1      1

$ ps
pid      ppid    name      state    size      start time      ticks    sched    cpu    thd cnt is thrd thrd #
1         1      init      sleep    12288     2020-04-30 10:43:22    9        43
2         1       sh        sleep    16384     2020-04-30 10:43:22    6        36
6         5      thtst2    runble   77824     2020-04-30 10:43:33   418      420
5         1      thtst2    runble   77824     2020-04-30 10:43:33    6       872      6      y      1
7         5      thtst2    run      77824     2020-04-30 10:43:33   419      420      2      y      2
8         5      thtst2    runble   77824     2020-04-30 10:43:33   417      418      y      3
9         5      thtst2    run      77824     2020-04-30 10:43:33   416      417      0      y      4
10        5      thtst2    run      77824     2020-04-30 10:43:33   416      416      3      y      5
11        5      thtst2    runble   77824     2020-04-30 10:43:33   415      415      y      6
15        2       ps        run      12288     2020-04-30 10:43:39    0         1      1
```